

2 System Specifications

System:

Processor: Intel Core i7-8550U CPU @ 1.80GHz 1.99 GHz

Installed RAM: 8.00 GB

System type: 64-bit operating system

Edition: Windows 10 Home

Disk Sequential 64.0 Read: 128.05 MB/s Disk Random 16.0 Read: 1.49 MB/s

PostgreSQL:

version 13.0

3 Query Optimization

3.1 Equivalent queries: Query-6

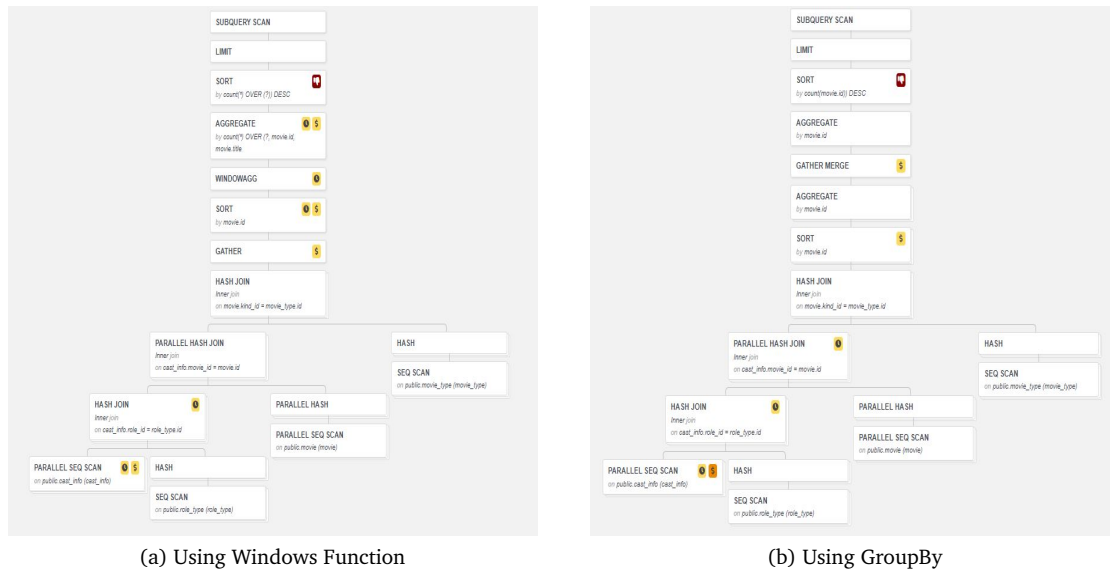


Figure 2: Query Plans for query 6

Node Type	Count	Time	
Aggregate	1	2s 451ms	21%
Hash Join	3	2s 344ms	20%
Seq Scan	4	2s 164ms	18%
Sort	2	1s 524ms	14%
WindowAgg	1	1s 388ms	12%
Gather	1	1s 138ms	10%
Hash	3	437ms	4%
Subquery Scan	1	0.004ms	0%
Limit	1	0.002ms	0%

(a) Using Windows Function

Node Type	Count	Time	
Hash Join	3	2s 188ms	37%
Seq Scan	4	1s 858ms	31%
Aggregate	2	646ms	11%
Sort	2	565ms	10%
Gather Merge	1	390ms	7%
Hash	3	237ms	4%
Subquery Scan	1	0.006ms	0%
Limit	1	0.004ms	0%

(b) Using GroupBy

Figure 3: Per Node Stats for query 6 during a particular hot run

Query-6 was originally written using Windows Function. The Aggregation in the query with the Windows Function is slower (see Figure 3) since it expands each row into the set of rows that represent the window associated with it, and then performs the aggregation. This leads to more reads and writes to the tempdb before being sorted and is thus by nature, slow.

It is possible to eliminate multiple read-writes. This can be done by an equivalent query using GroupBy can be written wherein the I/O operations happen once for Aggregation. This leads to better performance as can be seen from the results of the average run-time of hot-runs.

Table 1					
	Cold	Hot-1	Hot-2	Hot-3	Hot-Avg
Over	120s	12s	10.8s	10.6s	11.13s
GroupBy	119s	5.9s	7.7s	5.8s	6,46s

3.2 Indexes: Query-1

For Query-1, as can be seen from Figure 4, due to the nature of the query (WHERE clause filters specific rows on a very large table), the sequential scans to find the relevant records (person.name = 'Wayne, John' or person.name = 'Streep, Meryl') and similarly to form the joined table, an immediate performance improvement can be expected using indexes.

Therefore, in accordance to the WHERE clause of the query, we clustered person.name for the person table, and indexed on personInfo(personId, infoId). The improvement in query execution time is almost a 1000-fold as can be seen from the results in Table 2.

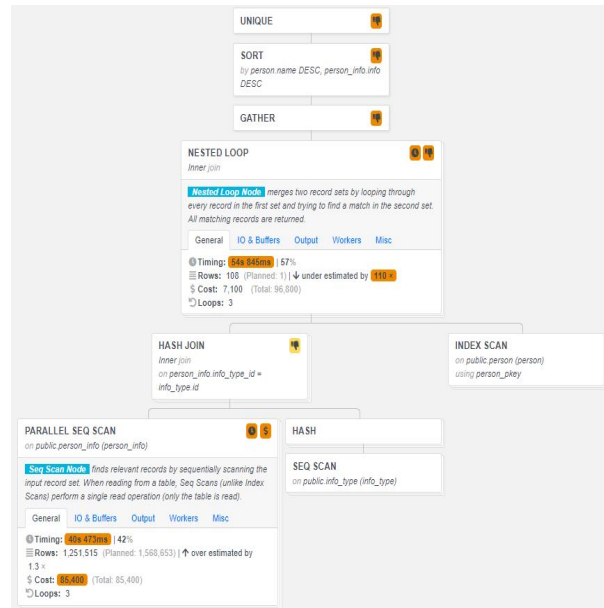


Figure 4: Query plan for query-1

Table 2					
	Cold	Hot-1	Hot-2	Hot-3	Hot-Avg
Normal	96s	1128ms	1130ms	1140ms	1132ms
Index	11ms	2.4ms	1.4ms	1.5ms	1.8ms

3.3 Materialized View - Query 4

Query-4 was originally written using INTERSECTION, which after analyzing the query plan, was found to have room for improvement (reads happening in 2 parts). The query was improved by using ALIASES which reduced the execution time by almost half. Since the nature of the query was to scan through thousands of rows but return a handful, creating a materialized view of the subquery used in Query-4 brought down the execution speed from 71.6s to 13.6ms (see Table 3).

Table 2					
	Cold	Hot-1	Hot-2	Hot-3	Hot-Avg
Intersection	3037s	210s	87s	59s	118.6s
Alias	170s	136s	43s	36s	71.6s
Material View	17.4	13.4ms	13.5ms	14ms	13.6ms

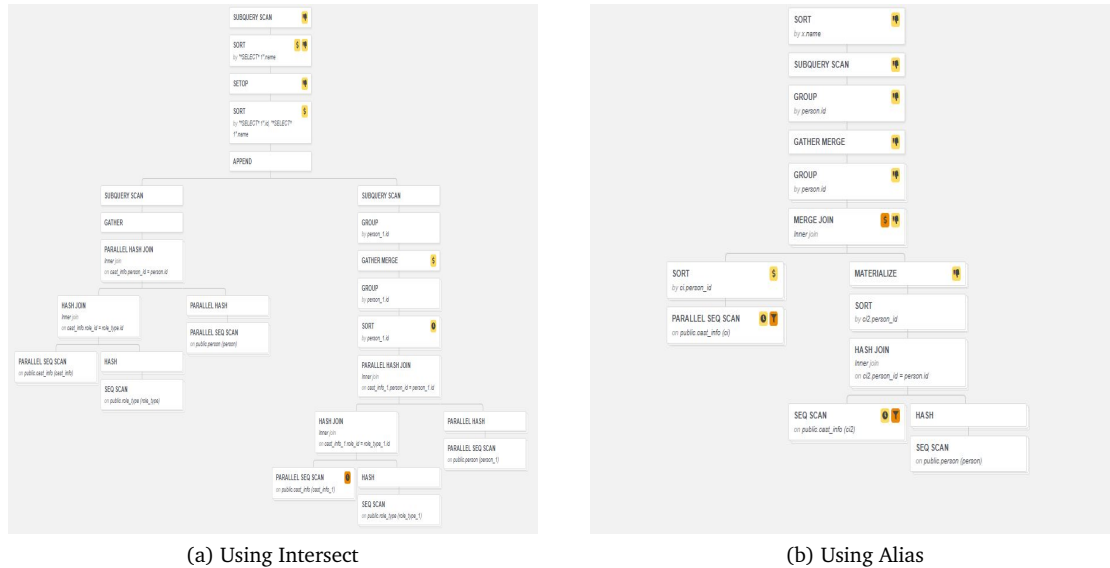


Figure 5: Query plans for query 4 during a particular hot run

4 Comparison to Pandas

We created a jupyter notebook in which we imported the whole database into pandas dataframes. Then we recreated Query-1 in python with pandas and measured the time.

Load-time of all tables = 342.194s

Load-time of required tables = 13.837s

Execution-time = 0.787s

Next we used pandas.readsqlquery after initializing the connection to postgresSQL to run Query-1.

Execution-time = 0.0518s

Thus we see that atleast for this particular query postgresSQL is almost 15x faster. It can be safely concluded that postgresSQL is way faster to use than pandas to achieve the same goal.

5 APPENDIX

Equivalent queries - Query 6

Original query with Over()

Cold-run: <https://explain.dalibo.com/plan/Uptv>

Hot-run1: <https://explain.dalibo.com/plan/c0>

Hot-run2: <https://explain.dalibo.com/plan/KVs>

Hot-run3: <https://explain.dalibo.com/plan/PRK>

Equivalent query using GroupBy

Cold-run: <https://explain.dalibo.com/plan/49c>

Hot-run 1: <https://explain.dalibo.com/plan/LCq>

Hot-run 2: <https://explain.dalibo.com/plan/axt>

Hot-run 3: <https://explain.dalibo.com/plan/XhD>

Indexes - Query 1

Original query

Cold-run: <https://explain.dalibo.com/plan/K8V>

Hot-run 1: <https://explain.dalibo.com/plan/gtT>

Hot-run 2: <https://explain.dalibo.com/plan/pX0>

Hot-run 3: <https://explain.dalibo.com/plan/3YW>

Query with clustered person table on name with ordered index and indexed on personInfo(personId, infoId)

Cold-run: <https://explain.dalibo.com/plan/xDD>

Hot-run 1: <https://explain.dalibo.com/plan/cua>

Hot-run 2: <https://explain.dalibo.com/plan/1BP>

Hot-run 3: <https://explain.dalibo.com/plan/sTC>

Materialized View - Query 6

Original query using Intersection

Cold-run: <https://explain.dalibo.com/plan/ddb>

Hot-run1: <https://explain.dalibo.com/plan/v1X>

Hot-run2: <https://explain.dalibo.com/plan/s93>

Hot-run3: <https://explain.dalibo.com/plan/AJk>

Equivalent query using aliases

Cold-run: <https://explain.dalibo.com/plan/eEB>

Hot-run1: <https://explain.dalibo.com/plan/6nu>

Hot-run2: <https://explain.dalibo.com/plan/hX8>

Hot-run3: <https://explain.dalibo.com/plan/mua>

Query using Materialized View

Cold-run: <https://explain.dalibo.com/plan/gpA>

Hot-run1: <https://explain.dalibo.com/plan/pN2>

Hot-run2: <https://explain.dalibo.com/plan/k52>

Hot-run3: <https://explain.dalibo.com/plan/OD2n>